class14R

Katelyn Brown A15891811

3/3/2022

Read our vaccination data.

Downloaded the most recently dated "Statewide COVID-19 Vaccines Administered by ZIP Code" CSV.

```
vax <- read.csv("covid19vaccinesbyzipcode_test.csv")
head(vax)
```

```
##   as_of_date zip_code_tabulation_area local_health_jurisdiction         county
## 1 2021-01-05                   92549                 Riverside      Riverside
## 2 2021-01-05                   92130                 San Diego      San Diego
## 3 2021-01-05                   92397           San Bernardino San Bernardino
## 4 2021-01-05                   94563             Contra Costa   Contra Costa
## 5 2021-01-05                   94519             Contra Costa   Contra Costa
## 6 2021-01-05                   91042              Los Angeles    Los Angeles
##   vaccine_equity_metric_quartile               vem_source
## 1                              3 Healthy Places Index Score
## 2                              4 Healthy Places Index Score
## 3                              3 Healthy Places Index Score
## 4                              4 Healthy Places Index Score
## 5                              3 Healthy Places Index Score
## 6                              2 Healthy Places Index Score
##   age12_plus_population age5_plus_population persons_fully_vaccinated
## 1                2348.4                 2461                        NA
## 2               46300.3                53102                        61
## 3                3695.6                 4225                        NA
## 4               17216.1                18896                        NA
## 5               16861.2                18678                        NA
## 6               23962.2                25741                        NA
##   persons_partially_vaccinated percent_of_population_fully_vaccinated
## 1                           NA                                    NA
## 2                           27                                0.001149
## 3                           NA                                    NA
## 4                           NA                                    NA
## 5                           NA                                    NA
## 6                           NA                                    NA
##   percent_of_population_partially_vaccinated
## 1                                         NA
## 2                                   0.000508
## 3                                         NA
## 4                                         NA
```

```
## 5                                        NA
## 6                                        NA
##   percent_of_population_with_1_plus_dose booster_recip_count
## 1                                     NA                  NA
## 2                               0.001657                  NA
## 3                                     NA                  NA
## 4                                     NA                  NA
## 5                                     NA                  NA
## 6                                     NA                  NA
##                                                          redacted
## 1 Information redacted in accordance with CA state privacy requirements
## 2 Information redacted in accordance with CA state privacy requirements
## 3 Information redacted in accordance with CA state privacy requirements
## 4 Information redacted in accordance with CA state privacy requirements
## 5 Information redacted in accordance with CA state privacy requirements
## 6 Information redacted in accordance with CA state privacy requirements
```

Q1. What column details the total number of people fully vaccinated?

The persons_fully_vaccinated column details the total number of people fully vaccinated.

Q2. What column details the Zip code tabulation area?

The zip_code_tabulation_area column details the Zip code tabulation area.

Q3. What is the earliest date in this dataset?

The earliest date in the dataset is 2021-01-05.

Q4. What is the latest date in the dataset?

```
vax$as_of_date[nrow(vax)]
```

```
## [1] "2022-03-01"
```

The latest date in the dataset is 2022-03-01.

```
# install.packages(skimr)
library(skimr)
skimr::skim(vax)
```

Table 1: Data summary

Name	vax
Number of rows	107604
Number of columns	15
Column type frequency:	
character	5

Table 1: Data summary

numeric	10
Group variables	None

Variable type: character

skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
as_of_date	0	1	10	10	0	61	0
local_health_jurisdiction	0	1	0	15	305	62	0
county	0	1	0	15	305	59	0
vem_source	0	1	15	26	0	3	0
redacted	0	1	2	69	0	2	0

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
zip_code_tabulation_area	0	1.00	93665.11	1817.39	90001	92257.75	93658.50	95380.50	97635.0	
vaccine_equity_metric_quartile	5307	0.95	2.44	1.11	1	1.00	2.00	3.00	4.0	
age12_plus_population	0	1.00	18895.04	18993.91	0	1346.95	13685.10	31756.12	88556.7	
age5_plus_population	0	1.00	20875.24	21106.02	0	1460.50	15364.00	34877.00	101902.0	
persons_fully_vaccinated	18338	0.83	12155.61	13063.88	11	1066.25	7374.50	20005.00	77744.0	
persons_partially_vaccinated	18338	0.83	831.74	1348.68	11	76.00	372.00	1076.00	34219.0	
percent_of_population_fully_vaccinated	18338	0.83	0.51	0.26	0	0.33	0.54	0.70	1.0	
percent_of_population_partially_vaccinated	18338	0.83	0.05	0.09	0	0.01	0.03	0.05	1.0	
percent_of_population_with_1_plus_dose	18338	0.83	0.54	0.28	0	0.36	0.58	0.75	1.0	
booster_recip_count	64317	0.40	4100.55	5900.21	11	176.00	1136.00	6154.50	50602.0	

Q5. How many numeric columns are in this dataset?

There are 9 numeric columns in the dataset.

Q6. Note that there are "missing values" in the dataset. How many NA values there in the persons_fully_vaccinated column?

There are 18338 NA values in the persons_fully_vaccinated column.

Q7. What percent of persons_fully_vaccinated values are missing (to 2 significant figures)?

17% of persons_fully_vaccinated is missing.

Q8. Why might this data be missing?

This data could be missing due to lack of reporting, or possibly lack of access. Some counties may also not be reporting vaccination rates.

```
# install.packages("lubridate")
library(lubridate)
```

```
##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
##
##     date, intersect, setdiff, union
```

```
today()
```

```
## [1] "2022-03-03"
```

```
age <- today() - ymd("2001-11-16")
age
```

```
## Time difference of 7412 days
```

```
time_length(age, "year")
```

```
## [1] 20.29295
```

```
vax$as_of_date <- ymd(vax$as_of_date)
today() - vax$as_of_date[1]
```

```
## Time difference of 422 days
```

Q9. How many days have passed since the last update of the dataset?

```
today() - vax$as_of_date[nrow(vax)]
```

```
## Time difference of 2 days
```

```
# Determine how many days the dataset spans
vax$as_of_date[nrow(vax)] - vax$as_of_date[1]
```

```
## Time difference of 420 days
```

2 days have passed since the last update of the dataset.

Q10. How many unique dates are in the dataset (i.e. how many different dates are detailed)?

```
length(unique(vax$as_of_date))
```

```
## [1] 61
```

There are 61 unique dates in the dataset.

Working with Zip Codes

First, download the zipcodeR package and load it in the library.

```
# install.packages("zipcodeR")
library(zipcodeR)
```

```
geocode_zip('92037')
```

```
## # A tibble: 1 x 3
##   zipcode   lat   lng
##   <chr>   <dbl> <dbl>
## 1 92037    32.8 -117.
```

```
zip_distance('92037', '92109')
```

```
##   zipcode_a zipcode_b distance
## 1     92037     92109     2.33
```

```
reverse_zipcode(c('92037', '92109'))
```

```
## # A tibble: 2 x 24
##   zipcode zipcode_type major_city post_office_city common_city_list county state
##   <chr>   <chr>        <chr>      <chr>                     <blob> <chr>  <chr>
## 1 92037   Standard     La Jolla   La Jolla, CA           <raw 20 B> San D~ CA
## 2 92109   Standard     San Diego  San Diego, CA          <raw 21 B> San D~ CA
## # ... with 17 more variables: lat <dbl>, lng <dbl>, timezone <chr>,
## #   radius_in_miles <dbl>, area_code_list <blob>, population <int>,
## #   population_density <dbl>, land_area_in_sqmi <dbl>,
## #   water_area_in_sqmi <dbl>, housing_units <int>,
## #   occupied_housing_units <int>, median_home_value <int>,
## #   median_household_income <int>, bounds_west <dbl>, bounds_east <dbl>,
## #   bounds_north <dbl>, bounds_south <dbl>
```

```
zipdata <- reverse_zipcode(vax$zip_code_tabulation_area)
zipdata
```

```
## # A tibble: 1,764 x 24
##    zipcode zipcode_type major_city  post_office_city common_city_list county
##    <chr>   <chr>        <chr>       <chr>                     <blob> <chr>
##  1 90001   Standard     Los Angeles Los Angeles, CA        <raw 44 B> Los Angel~
##  2 90002   Standard     Los Angeles Los Angeles, CA        <raw 47 B> Los Angel~
##  3 90003   Standard     Los Angeles Los Angeles, CA        <raw 23 B> Los Angel~
##  4 90004   Standard     Los Angeles Los Angeles, CA        <raw 34 B> Los Angel~
##  5 90005   Standard     Los Angeles Los Angeles, CA        <raw 34 B> Los Angel~
##  6 90006   Standard     Los Angeles Los Angeles, CA        <raw 23 B> Los Angel~
##  7 90007   Standard     Los Angeles Los Angeles, CA        <raw 37 B> Los Angel~
##  8 90008   Standard     Los Angeles Los Angeles, CA        <raw 53 B> Los Angel~
##  9 90010   Standard     Los Angeles Los Angeles, CA        <raw 23 B> Los Angel~
## 10 90011   Standard     Los Angeles Los Angeles, CA        <raw 23 B> Los Angel~
```

```
## # ... with 1,754 more rows, and 18 more variables: state <chr>, lat <dbl>,
## #   lng <dbl>, timezone <chr>, radius_in_miles <dbl>, area_code_list <blob>,
## #   population <int>, population_density <dbl>, land_area_in_sqmi <dbl>,
## #   water_area_in_sqmi <dbl>, housing_units <int>,
## #   occupied_housing_units <int>, median_home_value <int>,
## #   median_household_income <int>, bounds_west <dbl>, bounds_east <dbl>,
## #   bounds_north <dbl>, bounds_south <dbl>
```

Filter to only work with San Diego county.

```
# Base R
dim(vax[vax$county == "San Diego", ])
```

```
## [1] 6527   15
```

```
# Dplyr
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##     filter, lag

## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union
```

```
sd <- filter(vax, county == "San Diego")
dim(sd)
```

```
## [1] 6527   15
```

Q11. How many distinct zip codes are listed for San Diego County?

```
length(unique(sd$zip_code_tabulation_area))
```

```
## [1] 107
```

There are 107 unique zip codes in San Diego County.

Q12. What San Diego County Zip code area has the largest 12 + Population in this dataset?

```
sd[which.max(sd$age12_plus_population), "zip_code_tabulation_area"]
```

```
## [1] 92154
```

92154 has the largest 12+ population in the dataset.

Q13. What is the overall average "Percent of Population Fully Vaccinated" value for all San Diego "County" as of "2022-03-01"?

```
sd$as_of_date[nrow(sd)]
```

```
## [1] "2022-03-01"
```

```
# Filter to day
sd.latest <- filter(sd, as_of_date == "2022-03-01")
mean(sd.latest$percent_of_population_fully_vaccinated, na.rm = TRUE)*100
```

```
## [1] 70.52904
```

```
summary(sd.latest$percent_of_population_fully_vaccinated, na.rm = T)
```

```
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.    NA's
## 0.01017 0.65132 0.72452 0.70529 0.82567 1.00000       1
```

70.53% of San Diego county is fully vaccinated as of 2022-03-01.

> Q14. Using either ggplot or base R graphics make a summary figure that shows the distribution of Percent of Population Fully Vaccinated values as of "2022-03-01"?

```
library(ggplot2)
ggplot(sd.latest) + aes(x = sd.latest$percent_of_population_fully_vaccinated) + geom_histogram() + labs
```

```
## Warning: Use of `sd.latest$percent_of_population_fully_vaccinated` is
## discouraged. Use `percent_of_population_fully_vaccinated` instead.
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 1 rows containing non-finite values (stat_bin).
```

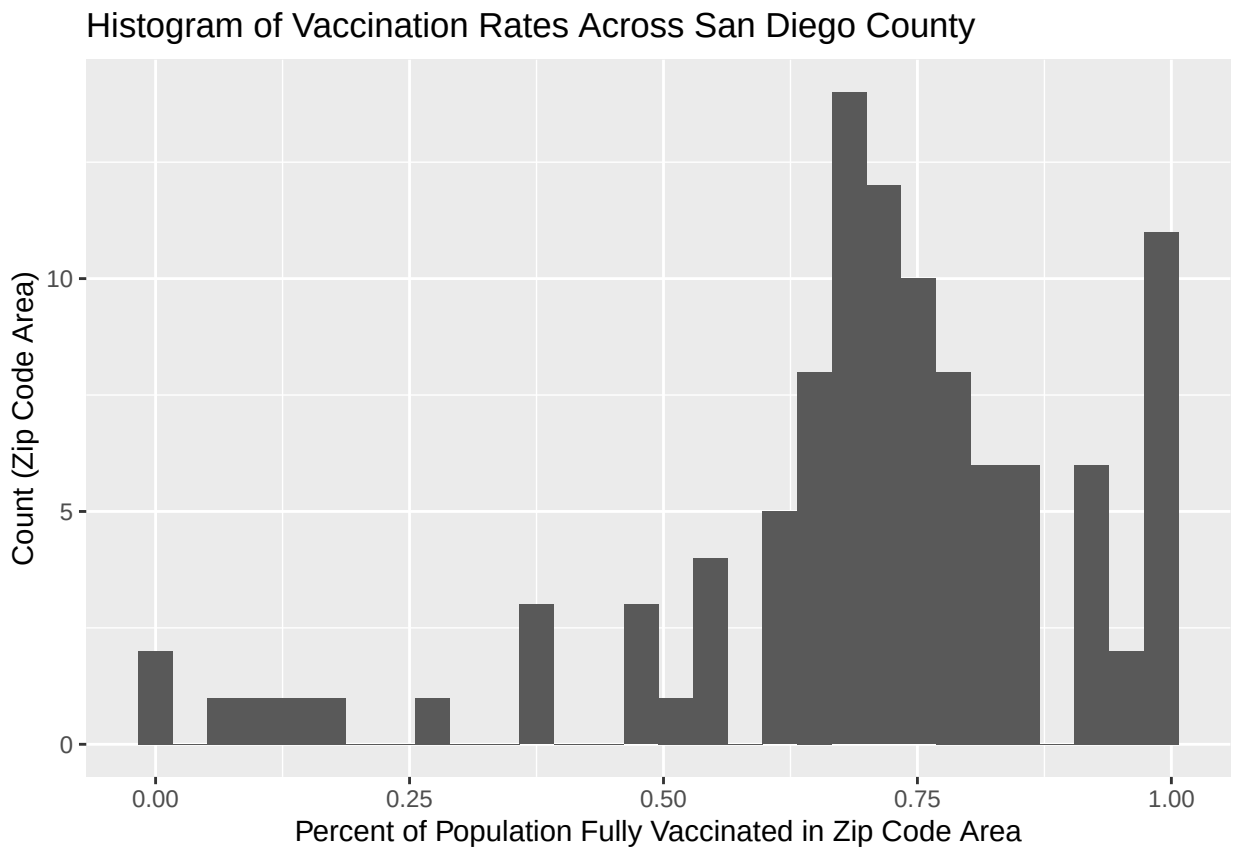

Filter to focus on UCSD.

```
ucsd <- filter(sd, zip_code_tabulation_area=="92037")
ucsd[1,]$age5_plus_population
```

```
## [1] 36144
```

Q15. Using ggplot make a graph of the vaccination rate time course for the 92037 ZIP code area.

```
baseplot <- ggplot(ucsd) +
  aes(x = as_of_date, y = percent_of_population_fully_vaccinated) +
  geom_point() +
  geom_line(group=1) + theme_linedraw() +
  ylim(c(0,1)) +
  labs(x = "Date", y="Percent Vaccinated", title = "Vaccination Rate in La Jolla, CA 92037")
baseplot
```

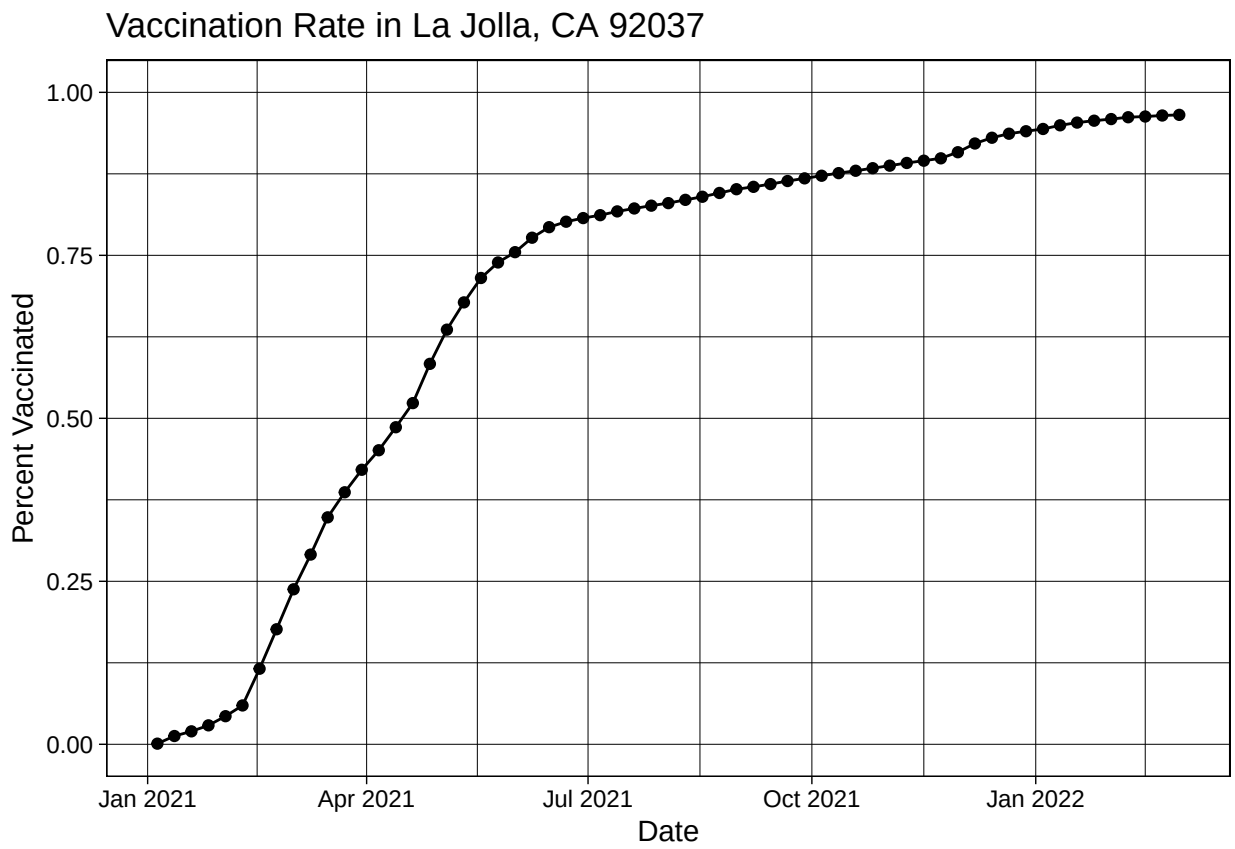

Q16. Calculate the mean "Percent of Population Fully Vaccinated" for ZIP code areas with a population as large as 92037 (La Jolla) as_of_date "2022-03-01". Add this as a straight horizontal line to your plot from above with the geom_hline() function?

```
# Subset other CA zip codes with populations as big as 92037
vax.36 <- filter(vax, age5_plus_population > 36144 & as_of_date == "2022-03-01")
head(vax.36)
```

```
##   as_of_date zip_code_tabulation_area local_health_jurisdiction      county
## 1 2022-03-01                   95628                Sacramento  Sacramento
## 2 2022-03-01                   90808                Long Beach Los Angeles
## 3 2022-03-01                   92507                 Riverside   Riverside
## 4 2022-03-01                   92626                    Orange      Orange
## 5 2022-03-01                   93257                    Tulare      Tulare
## 6 2022-03-01                   90011               Los Angeles Los Angeles
##   vaccine_equity_metric_quartile                vem_source
## 1                              3 Healthy Places Index Score
## 2                              4 Healthy Places Index Score
## 3                              1 Healthy Places Index Score
## 4                              3 Healthy Places Index Score
## 5                              1 Healthy Places Index Score
## 6                              1 Healthy Places Index Score
##   age12_plus_population age5_plus_population persons_fully_vaccinated
## 1               35579.0                38694                    28842
## 2               33952.3                37179                    29383
```

```
## 3               51432.5               55253                    34455
## 4               44238.8               47883                    33767
## 5               61519.8               70784                    42919
## 6               87902.8              101902                    65342
##   persons_partially_vaccinated percent_of_population_fully_vaccinated
## 1                         1990                                0.745387
## 2                         2112                                0.790312
## 3                         3947                                0.623586
## 4                         2937                                0.705198
## 5                         5868                                0.606338
## 6                        15255                                0.641224
##   percent_of_population_partially_vaccinated
## 1                                   0.051429
## 2                                   0.056806
## 3                                   0.071435
## 4                                   0.061337
## 5                                   0.082900
## 6                                   0.149703
##   percent_of_population_with_1_plus_dose booster_recip_count redacted
## 1                               0.796816               16913       No
## 2                               0.847118               17253       No
## 3                               0.695021               15073       No
## 4                               0.766535               17595       No
## 5                               0.689238               17740       No
## 6                               0.790927               19928       No
```

```
# Find mean of this data
vax.36mean <- mean(vax.36$percent_of_population_fully_vaccinated, na.rm = TRUE)
vax.36mean
```

```
## [1] 0.7353974
```

```
# Add mean to baseplot
baseplot + geom_hline(yintercept = vax.36mean, linetype = 2, alpha = 0.5, color = "red")
```

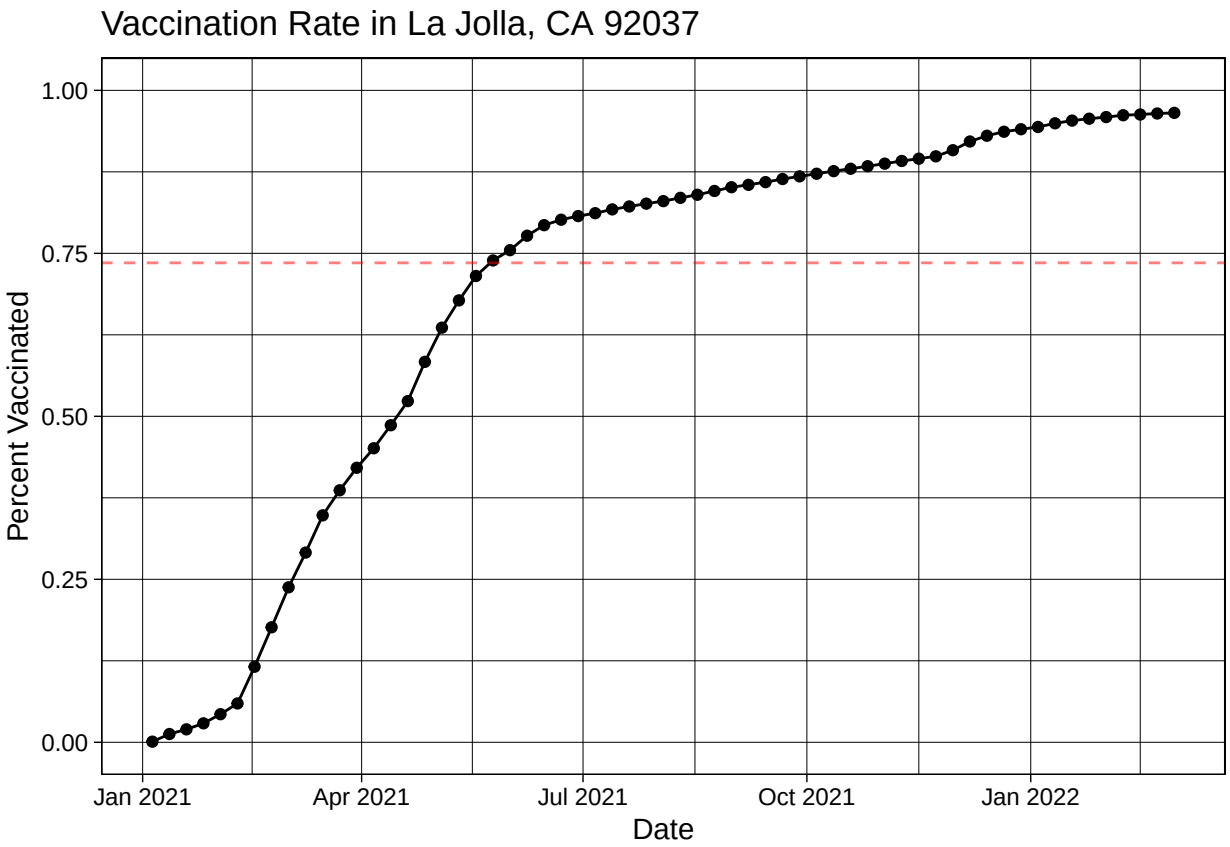

Q17. What is the 6 number summary (Min, 1st Qu., Median, Mean, 3rd Qu., and Max) of the "Percent of Population Fully Vaccinated" values for ZIP code areas with a population as large as 92037 (La Jolla) as_of_date "2022-03-01"?

```
summary(vax.36$percent_of_population_fully_vaccinated)
```

```
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.3890  0.6554  0.7350  0.7354  0.8044  1.0000
```

Q18. Using ggplot generate a histogram of this data.

```
ggplot(vax.36) + aes(percent_of_population_fully_vaccinated) + geom_histogram() + xlim(c(0,1))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 2 rows containing missing values (geom_bar).
```

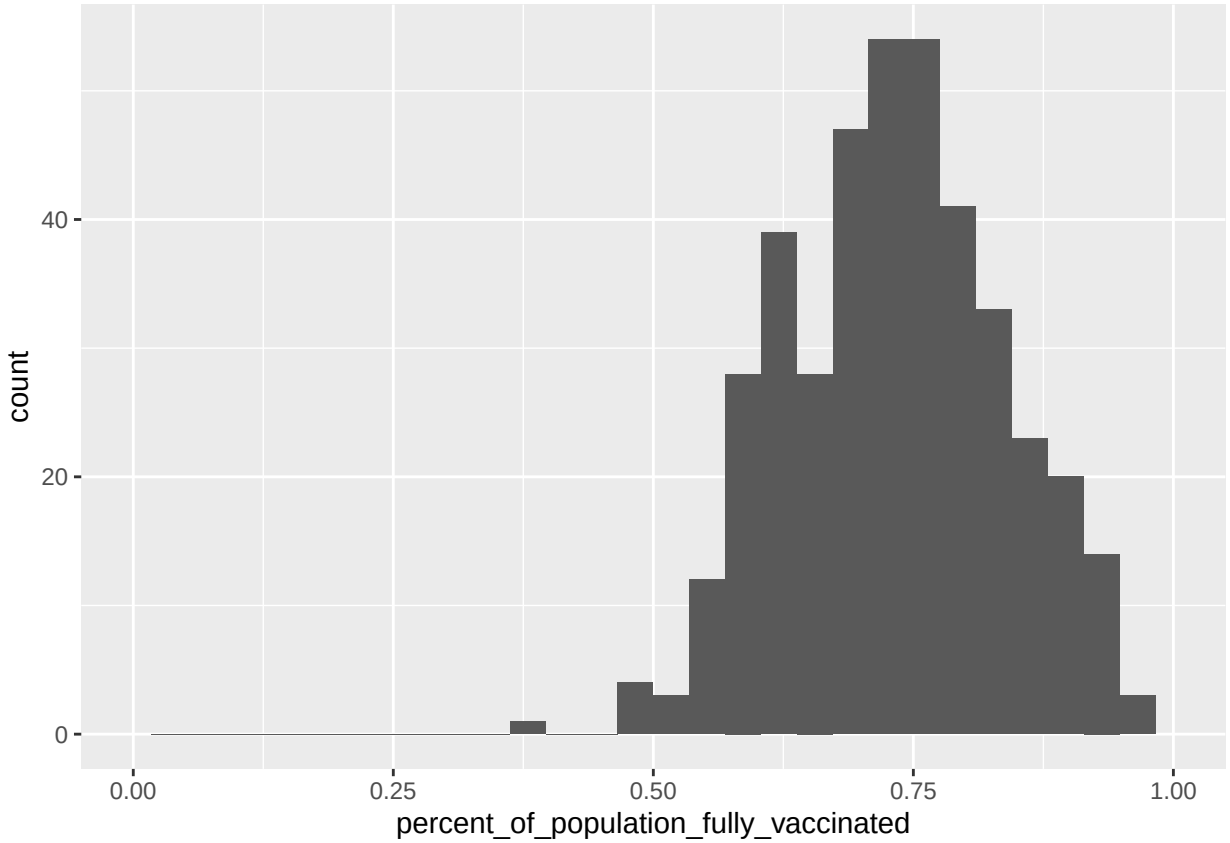

Q19. Is the 92109 and 92040 ZIP code areas above or below the average value you calculated for all these above?

```
vax %>% filter(as_of_date == "2022-03-01") %>%
  filter(zip_code_tabulation_area=="92109") %>%
  select(percent_of_population_fully_vaccinated)
```

```
##   percent_of_population_fully_vaccinated
## 1                               0.723778
```

```
vax %>% filter(as_of_date == "2022-03-01") %>%
  filter(zip_code_tabulation_area=="92040") %>%
  select(percent_of_population_fully_vaccinated)
```

```
##   percent_of_population_fully_vaccinated
## 1                               0.551981
```

92109 and 92040 zip codes are below the average of 0.735 that was calculated previously.

Q20. Finally make a time course plot of vaccination progress for all areas in the full dataset with a age5_plus_population > 36144.

```
vax.36.all <- filter(vax, age5_plus_population > 36144)
ggplot(vax.36.all) +
  aes(as_of_date,
      percent_of_population_fully_vaccinated,
      group = zip_code_tabulation_area) +
  geom_line(alpha = 0.2, color = "blue") +
  ylim(0, 1) +
  labs(x = "Date", y = "Percent Vaccinated",
       title = "Vaccination Rate Across California",
       subtitle = "Only areas with populations above 36k are shown.") +
  geom_hline(yintercept = vax.36mean, linetype = 2)
```

```
## Warning: Removed 311 row(s) containing missing values (geom_path).
```